

The background of the slide features a large, faint watermark of the University of Padua seal. The seal is circular, with the Latin text "UNIVERSITAS STUDII PADUENSIS" around the perimeter and "MCCXXII" at the bottom. In the center is a shield depicting two figures, one seated and one standing, with a wheel and stars above them.

# Deep Learning

*LM Computer Science & Data Science*

*2<sup>nd</sup> semester - 6 CFU*

*Nicolò Navarin & Alessandro Sperduti*

# Probability / Information theory

A primer on probability and information theory (chapter 3)

Maximum Likelihood estimation (section 5.5)

# Probability

- **Random variable:** a variable that can take different values randomly
- Example: Tossing a coin: we could get Heads or Tails.
  - Heads=0 and Tails=1
  - In each experiment, Random Variable  $x$  can be either 0 or 1

*Random  
Variable*

*Possible  
Values*


*Random  
Events*

$$X = \begin{cases} 0 \\ 1 \end{cases}$$



# Probability Distributions

- **Probability Distribution:** A description of how likely a random variable  $x$  (or a set of random variables) is to take each of its possible states
- Discrete variables  $\rightarrow$  Probability Mass Function
  - Domain of  $P$  is the set of all possible states of  $x$  ( $k$  different values)
  - $\forall x \in \mathcal{X} \ 0 \leq P(x = x) \leq 1$
  - $\sum_{x \in \mathcal{X}} P(x) = 1$
- E.g. Uniform distribution  $\forall_{x \in \mathcal{X}} P(x = x) = \frac{1}{k}$



# Probability Distributions

- **Joint probability:** probability distribution over many variables  $P(x = x, y = y)$  or  $P(x, y)$

- **Marginalization (sum rule):**

$$\forall x \in \mathbf{x} \ P(x = x) = \sum_y P(x = x, y = y)$$


- When dealing with **continuous variables**, a Probability distribution is described by a Probability Density Function (PDF)
  - Domain of  $p$  is the set of all possible states of  $x$
  - $\forall x \in \mathbf{x} \ P(x) \geq 0$
  - $\int p(x) dx = 1$
- E.g. Gaussian distribution

# Gaussian distribution

$$\mathcal{N}(x; \mu, \sigma^2) = \sqrt{\frac{1}{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2} (x - \mu)^2\right)$$

Parametrized by

PDF of  
Gaussian  
distribution




# Entropy

- Shannon Entropy (discrete variable)

$$H(x) = -\mathbb{E}_{x \sim P(x)} [\log P(x)]$$

- Expected amount of (self-)information in an event drawn from distribution  $P$
- Lower bound on the number of bits needed on average to encode a symbol drawn from that distribution

Entropy  $H(X)$  of a coin flip, measured in bits, graphed versus the bias of the coin  $P(x = 1)$ , where  $x = 1$  represents a result of heads.



# Kullback-Leibler divergence and Cross Entropy

- Let's consider two probability distributions  $P(x)$  and  $Q(x)$
- Measure how different they are

$$D_{KL}(P \parallel Q) = \mathbb{E}_{x \sim P} \left[ \log \frac{P(x)}{Q(x)} \right] = \mathbb{E}_{x \sim P} [\log P(x) - \log Q(x)]$$

- It is not a true distance because it is not symmetric

- Cross Entropy

$$H(P, Q) = H(P) + D_{KL}(P \parallel Q) = - \mathbb{E}_{x \sim P} [\log Q(x)]$$

Note: minimizing CE of  $P$  w.r.t.  $Q$  is equivalent to minimize KL divergence between  $P$  and  $Q$  (if  $P$  is given,  $H(P)$  and  $\mathbb{E}_{x \sim P} [\log P(x)]$  are constants)



# Maximum likelihood estimation

- Principled way to derive estimators (models)
- Consider  $n$  examples  $Tr = \{\mathbf{x}^1, \dots, \mathbf{x}^n\}$  drawn i.i.d. from  $p_{data}(\mathbf{x})$
- Let us consider a family of parametric probability distributions (models)  $p_{model}(\mathbf{x}; \boldsymbol{\theta})$ .
  - $p_{model}(\mathbf{x}; \boldsymbol{\theta})$  maps a point  $\mathbf{x}$  to a real number estimating  $p_{data}(\mathbf{x})$
  - Maximum Likelihood estimation for  $\boldsymbol{\theta}$  is

$$\boldsymbol{\theta}_{ML} = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} p_{model}(Tr; \boldsymbol{\theta}) = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} \prod_{i=1}^n p_{model}(\mathbf{x}^{(i)}; \boldsymbol{\theta})$$

Assumption: independence

# ..A side note on maximum likelihood

- ML is a special case of maximum a posteriori estimation (MAP) that assumes a uniform prior distribution
- MAP and maximum likelihood approach makes predictions using a single point estimate of  $\theta$
- the Bayesian approach is to make predictions using a full probability distribution over  $\theta$

# ..A side note on maximum likelihood

Given a new instance  $x$ , what is the most probable *classification* ?

- ▶  $h_{MAP}(x)$ , in general, is not the most probable classification!

Example: let's consider:

- ▶ three possible hypotheses:

$$P(h_1|D) = .4, \quad P(h_2|D) = .3, \quad P(h_3|D) = .3$$

- ▶ given a new instance  $x$ ,

$$h_1(x) = +, \quad h_2(x) = -, \quad h_3(x) = -$$

- ▶ what is the most probable classification for  $x$ ?

Bayes optimal classifier! (not covered in this course)

# Maximum likelihood estimation

- Taking the product of many probabilities is numerically unstable. We can apply the log and the arg max does not change

$$\boldsymbol{\theta}_{ML} = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} \sum_{i=1}^n \log p_{model}(\mathbf{x}^{(i)}; \boldsymbol{\theta})$$

- We can equivalently divide by  $n$  to express ML as an expectation over training data

$$\boldsymbol{\theta}_{ML} = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} \mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log p_{model}(\mathbf{x}; \boldsymbol{\theta})]$$

- ML minimizes the dissimilarity between  $\hat{p}_{data}$  and  $p_{model}$ , measured by the KL divergence (actually cross entropy, see next slide)

# ML estimation as KL divergence

$$D_{KL}(\hat{p}_{data} \parallel p_{model}) = \mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log \hat{p}_{data}(\mathbf{x}) - \log p_{model}(\mathbf{x}; \boldsymbol{\theta})]$$

- The term on the left does not depend on the model.
- To minimize the KL, we need only to minimize
$$\arg \min_{\boldsymbol{\theta}} -\mathbb{E}_{\mathbf{x} \sim \hat{p}_{data}} [\log p_{model}(\mathbf{x}; \boldsymbol{\theta})]$$
- That is the same equation of ML in previous slide
- It also corresponds to minimizing the **cross-entropy** between the two distributions (5 slides back)

# Conditional Probability

- Probability of an event, given that some other event has happened.

$$P(a, b) = P(a|b)P(b) \quad \text{Chain rule of probability}$$

The diagram illustrates Bayes' rule with the following components:

- LIKELIHOOD** (orange text): the probability of "B" being TRUE given that "A" is TRUE. An arrow points from this text to the  $P(B|A)$  term in the numerator.
- PRIOR** (teal text): the probability of "A" being TRUE. An arrow points from this text to the  $P(A)$  term in the numerator.
- POSTERIOR** (green text): the probability of "A" being TRUE given that "B" is TRUE. An arrow points from this text to the  $P(A|B)$  term on the left side of the equation.
- The probability of "B" being TRUE** (pink text): An arrow points from this text to the  $P(B)$  term in the denominator.

The equation is presented as:

$$P(A|B) = \frac{P(B|A) P(A)}{P(B)}$$

Bayes' rule

# Conditional log likelihood

- We can use ML to estimate a **conditional** probability  $P(\mathbf{y}|\mathbf{x}; \boldsymbol{\theta})$  to predict  $\mathbf{y}$  given  $\mathbf{x}$  (supervised learning)

$$\boldsymbol{\theta}_{ML} = \arg \max_{\boldsymbol{\theta}} P(\mathbf{Y} | \mathbf{X}; \boldsymbol{\theta})$$

- If examples are i.i.d.

$$\boldsymbol{\theta}_{ML} = \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^n \log P(\mathbf{y}^{(i)} | \mathbf{x}^{(i)}; \boldsymbol{\theta})$$

# Why MSE? Linear regression as ML

- Let's think about the model as producing a conditional distribution  $p(y | \mathbf{x})$

Model prediction

- We define  $p(y | \mathbf{x}) = \mathcal{N}(y; \hat{y}(\mathbf{x}; \boldsymbol{\theta}), \sigma^2)$
- Our model produces  $\hat{y}(\mathbf{x}; \boldsymbol{\theta})$ , the **mean of a Gaussian distribution**

- For i.i.d. examples,

$$\begin{aligned}\boldsymbol{\theta}_{ML} &= \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^n \log p(y^{(i)} | \mathbf{x}^{(i)}; \boldsymbol{\theta}) \\ &= \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^n \log \left( \sqrt{\frac{1}{2\pi\sigma^2}} \exp \left( -\frac{1}{2\sigma^2} (y^{(i)} - \hat{y}^{(i)})^2 \right) \right)\end{aligned}$$



# Warning!

- We will use logarithm properties. Check Algebra cheat sheet if some of the rules applied in the next slide are not clear!

# Algebra Cheat Sheet

## Basic Properties & Facts

### Arithmetic Operations

$$ab + ac = a(b + c)$$

$$\left(\frac{a}{b}\right) \div \frac{c}{d} = \frac{a}{bc}$$

$$\frac{a}{b} + \frac{c}{d} = \frac{ad + bc}{bd}$$

$$\frac{a-b}{c-d} = \frac{b-a}{d-c}$$

$$\frac{ab+ac}{a} = b+c, \quad a \neq 0$$

### Exponent Properties

$$a^n a^m = a^{n+m}$$

$$(a^n)^m = a^{nm}$$

$$(ab)^n = a^n b^n$$

$$a^{-n} = \frac{1}{a^n}$$

$$\left(\frac{a}{b}\right)^{-n} = \left(\frac{b}{a}\right)^n = \frac{b^n}{a^n}$$

### Properties of Radicals

$$\sqrt[n]{a} = a^{\frac{1}{n}} \quad \sqrt[n]{ab} = \sqrt[n]{a} \sqrt[n]{b}$$

$$\sqrt[m]{\sqrt[n]{a}} = \sqrt[mn]{a} \quad \sqrt[n]{\frac{a}{b}} = \frac{\sqrt[n]{a}}{\sqrt[n]{b}}$$

$$\sqrt[n]{a^n} = a, \text{ if } n \text{ is odd}$$

$$\sqrt[n]{a^n} = |a|, \text{ if } n \text{ is even}$$

### Properties of Inequalities

If  $a < b$  then  $a + c < b + c$  and  $a - c < b - c$

If  $a < b$  and  $c > 0$  then  $ac < bc$  and  $\frac{a}{c} < \frac{b}{c}$

If  $a < b$  and  $c < 0$  then  $ac > bc$  and  $\frac{a}{c} > \frac{b}{c}$

### Properties of Absolute Value

$$|a| = \begin{cases} a & \text{if } a \geq 0 \\ -a & \text{if } a < 0 \end{cases}$$

$$|a| \geq 0 \quad |-a| = |a|$$

$$|ab| = |a||b| \quad \left|\frac{a}{b}\right| = \frac{|a|}{|b|}$$

$$|a+b| \leq |a| + |b| \quad \text{Triangle Inequality}$$

### Distance Formula

If  $P_1 = (x_1, y_1)$  and  $P_2 = (x_2, y_2)$  are two points the distance between them is

$$d(P_1, P_2) = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

### Complex Numbers

$$i = \sqrt{-1} \quad i^2 = -1 \quad \sqrt{-a} = i\sqrt{a}, \quad a \geq 0$$

$$(a+bi) + (c+di) = a+c + (b+d)i$$

$$(a+bi) - (c+di) = a-c + (b-d)i$$

$$(a+bi)(c+di) = ac - bd + (ad+bc)i$$

$$(a+bi)(a-bi) = a^2 + b^2$$

$$|a+bi| = \sqrt{a^2 + b^2} \quad \text{Complex Modulus}$$

$$\overline{(a+bi)} = a-bi \quad \text{Complex Conjugate}$$

$$(a+bi)(a+bi) = |a+bi|^2$$

## Logarithms and Log Properties

### Definition

$y = \log_b x$  is equivalent to  $x = b^y$

### Example

$$\log_5 125 = 3 \quad \text{because } 5^3 = 125$$

### Special Logarithms

$\ln x = \log_e x$  natural log

$\log x = \log_{10} x$  common log

where  $e = 2.718281828K$

### Logarithm Properties

$$\log_b b = 1 \quad \log_b 1 = 0$$

$$\log_b b^x = x \quad b^{\log_b x} = x$$

$$\log_b (x^r) = r \log_b x$$

$$\log_b (xy) = \log_b x + \log_b y$$

$$\log_b \left(\frac{x}{y}\right) = \log_b x - \log_b y$$

The domain of  $\log_b x$  is  $x > 0$

## Factoring and Solving

### Factoring Formulas

$$x^2 - a^2 = (x+a)(x-a)$$

$$x^2 + 2ax + a^2 = (x+a)^2$$

$$x^2 - 2ax + a^2 = (x-a)^2$$

$$x^2 + (a+b)x + ab = (x+a)(x+b)$$

$$x^3 + 3ax^2 + 3a^2x + a^3 = (x+a)^3$$

$$x^3 - 3ax^2 + 3a^2x - a^3 = (x-a)^3$$

$$x^3 + a^3 = (x+a)(x^2 - ax + a^2)$$

$$x^3 - a^3 = (x-a)(x^2 + ax + a^2)$$

$$x^{2n} - a^{2n} = (x^n - a^n)(x^n + a^n)$$

If  $n$  is odd then,

$$x^n - a^n = (x-a)(x^{n-1} + ax^{n-2} + \dots + a^{n-1})$$

$$x^n + a^n = (x+a)(x^{n-1} - ax^{n-2} + a^2x^{n-3} - \dots + a^{n-1})$$

### Completing the Square

$$\text{Solve } 2x^2 - 6x - 10 = 0$$

(1) Divide by the coefficient of the  $x^2$

$$x^2 - 3x - 5 = 0$$

(2) Move the constant to the other side.

$$x^2 - 3x = 5$$

(3) Take half the coefficient of  $x$ , square it and add it to both sides

$$x^2 - 3x + \left(-\frac{3}{2}\right)^2 = 5 + \left(-\frac{3}{2}\right)^2 = 5 + \frac{9}{4} = \frac{29}{4}$$

(4) Factor the left side

$$\left(x - \frac{3}{2}\right)^2 = \frac{29}{4}$$

(5) Use Square Root Property

$$x - \frac{3}{2} = \pm \sqrt{\frac{29}{4}} = \pm \frac{\sqrt{29}}{2}$$

(6) Solve for  $x$

$$x = \frac{3}{2} \pm \frac{\sqrt{29}}{2}$$

### Quadratic Formula

Solve  $ax^2 + bx + c = 0, \quad a \neq 0$

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

If  $b^2 - 4ac > 0$  - Two real unequal solns.

If  $b^2 - 4ac = 0$  - Repeated real solution.

If  $b^2 - 4ac < 0$  - Two complex solutions.

### Square Root Property

If  $x^2 = p$  then  $x = \pm \sqrt{p}$

### Absolute Value Equations/Inequalities

If  $b$  is a positive number

$$|p| = b \Rightarrow p = -b \quad \text{or} \quad p = b$$

$$|p| < b \Rightarrow -b < p < b$$

$$|p| > b \Rightarrow p < -b \quad \text{or} \quad p > b$$

# Why MSE? Linear regression as ML

Log product rule

$$\theta_{ML} = \arg \max_{\theta} \sum_{i=1}^n \log \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) + \log \exp \left( -\frac{1}{2\sigma^2} (y^{(i)} - \hat{y}^{(i)})^2 \right)$$

Log quotient rule

$$= \arg \max_{\theta} \sum_{i=1}^n \log(1) - \log(\sqrt{2\pi\sigma^2}) + \log \exp \left( -\frac{1}{2\sigma^2} (y^{(i)} - \hat{y}^{(i)})^2 \right)$$

Log power rule

$$= \arg \max_{\theta} \sum_{i=1}^n \log(1) - \log(\sqrt{2\pi\sigma^2}) - \frac{1}{2\sigma^2} (y^{(i)} - \hat{y}^{(i)})^2 \log(e)$$

Natural logarithm and Log power rule

$$= \arg \max_{\theta} \sum_{i=1}^n -\frac{1}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} (y^{(i)} - \hat{y}^{(i)})^2$$

Algebra

$$= \arg \max_{\theta} -\frac{n}{2} \log(2\pi\sigma^2) + \sum_{i=1}^n -\frac{1}{2} \left( \frac{(y^{(i)} - \hat{y}^{(i)})^2}{\sigma^2} \right)$$

$$= \arg \max_{\theta} -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

# Comparing ML with MSE

$$\theta_{ML} = \arg \max_{\theta} \underbrace{-\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2}_{\text{Does not depend on } \theta}$$

$$\theta_{MSE} = \arg \min_{\theta} \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

- The two functions give the same  $\theta$ !
- ML estimator is, asymptotically in the number of examples, the best (single) estimator.
- With a small number of examples, **regularization** strategies to reduce the variance (dedicated chapter).