

Deep Learning

LM Computer Science, Data Science, Cybersecurity
2nd semester - 6 CFU

References:

[Deep Learning Book](#) (main book)

[Mitchell](#) (machine learning concepts)

[Bishop](#) (machine learning)



Practical Methodology (chapter 11), i.e. How to use deep learning


General methodology

- Determine your goals
 - error metrics, target value
- Build an end-to-end pipeline, including:
 - data ingestion, model training
 - and the estimation of the performance metrics
- Be prepared to analyse your system, to understand:
 - which component is performing worse than expected
 - if poor performance is due to underfitting, overfitting or a bug
- Repeatedly make incremental changes (one at a time!)
 - E.g. gathering new data, changing hyperparameters, changing algorithms based on specific observations
 - Understand the effect of each change

Performance metrics

- Decide which error metric to use
 - This metric will drive all future decisions
 - Different from loss: e.g. error rate
 - Often different errors should be weighted differently (e.g. spam classification)
- Detection of rare events:
 - Consider a disease that only one in 1M people have
 - 99.9999% accuracy simply predicting that nobody has the disease
 - Precision: fraction of detections that are correct
 - Recall: fraction of true events that are detected
 - If the classifier outputs a score, metrics depend on a threshold

Precision/Recall curve: varying the threshold



Ex. ID	Prediction	Target
5	0.99	1
10	0.97	1
7	0.96	0
1	0.86	1
	...	
42	0.1	1
9	0.05	0

- Single number summarising the performance (without fixing the threshold): Area under Precision/Recall curve
- If we fix the threshold:

$$Fscore = \frac{2 \cdot precision \cdot recall}{precision + recall}$$

Desired level of performance

- Decide a desired level of performance
 - Based on the application (e.g. minimum accuracy for the system to be useful/trusted/safe in practice)
 - Previously published results
 - Zero error is often not possible!
 - Bayes error rate: lowest possible error rate for any classifier (seen in Machine Learning courses, or check it online)
 - May be limited by finite training data

Baseline models

- Baselines are simple Neural networks chosen to provide a first estimation of the achievable level of performance
- Depending on the complexity of the problem, you may start with machine learning (not deep learning) algorithms
- Define a baseline NN exploiting your knowledge:
 - if your data has a structure, exploit it (e.g. CNNs on images, RNNs on sequence data)
 - Use piecewise linear units (e.g. ReLU)
 - Chose e relatively simple architecture
 - Choose a standard learning algorithm: SGD with momentum and decaying learning rate, or Adam
 - Include regularisation (e.g. Dropout or weight decay)

Search for similar tasks

- If the task is similar to another task that has been studied:
 - Use the same model and algorithms that have been used for the other task
 - Use a pre-trained model

Determine whether to gather more data

- After training the first models, measure its performance and determine how to improve it
- Find out if the training performance is acceptable
 - Low training performance -> you are not exploiting the available data well. Increase model capacity: more layers or more units (Hyperparameters)
 - At the extreme, your data may be too noisy or not informative enough (or there are problems in data: check by yourself)
 - Re-collect the data measuring more variables or fixing the problems
- Measure test performance
 - If it is good, you are done
 - If it is not, try to decrease model capacity
 - Reduce the size of the model (Hyperparameters)
 - Add regularisation
 - If the generalisation gap is still high, gather more data if possible

Defects in data


- Can a human process it?



26624




Model capacity




Hyperparameter selection

- Deep learning models are sensitive to hyperparameters
- Learning rate can be tuned on the training set
 - All the others on the validation set
- Usually the performance follows a U-shaped curve for each hyperparameter
 - Optimal model capacity lies between underfitting and overfitting zones
 - Usually, high-capacity models well regularized perform best
- Intuition should drive the search for hyperparameters




Learning rate



Automatic hyper-parameter selection

- You just select the range for each hyperparameter
- Grid search: Test many different pre-defined options
- Random search: usually just some of the hyperparameters have a significant impact in performance
 - It can be more efficient than grid search
- Other approaches: e.g. model-based (Bayesian)



Debug your network

- Visualise the outputs of the model
 - E.g. object detection/classification, speech synthesis
- Visualise the worst mistakes
- If training error is high, try to overfit a tiny dataset
- Check the gradient (exploding or vanishing)
 - If you implemented the derivatives, compare them with numerical derivatives
- Check the activations: how many neurons fire?
 - ReLU can induce “dead” neurons

Example

- In the book (chapter 11), an example of multi-digit number recognition (from Google) is described, check it out.